

Given an exec's name and time window, **/exec-brief** autonomously discovers their Salesforce org, scores and prioritizes accounts, runs six parallel intelligence streams (M365, Slack, Salesforce, PowerBI, LinkedIn, public news), and delivers a cited, fact-backed brief — calendar prep cards, verified struggles, consumption health, pipeline, and ranked action items. Every finding requires a source citation. Nothing is fabricated.

USAGE

```
/exec-brief "Nishant Rama" 2w --html → 2-week branded HTML report saved to current dir
/exec-brief "Dom Miller" 3w --html focus on at-risk renewals → renewal-weighted brief; focus injected into every stream + synthesis
/exec-brief "Sarah Chen" last month --pdf → retrospective PDF, sanitized for distribution

Formats: --text Markdown to dir (default) · --html DT-branded dark report with sidebar nav · --pdf sanitized for distribution · --
resynthesize rerun synthesis without repulling sources
```

EXECUTION PHASES

PHASE	NAME	WHAT HAPPENS
Phase -1	MCP Pre-flight	Checks every required integration is present and OAuth-authenticated before any work begins. Blocks with exact install instructions if a required MCP is missing. LinkedIn absent → WebSearch/WebFetch fallback, skill continues without blocking.
Phase 0	Parse & Init	Resolves exec name, symmetric window (equal lookback for signals + lookahead for meetings), output format, and optional focus context. Focus context is injected verbatim into every subagent prompt and synthesis — it re-weights coverage without replacing any section.
Phase 1	Org Discovery	Resolves exec's Salesforce User ID, descends the org hierarchy (RDs → AEs, 2 levels, using SF ManagerId + SuccessFactors HR truth), maps all Customer and Prospect accounts, runs a completeness cross-check, then scores and caps focus accounts at 50. Meeting accounts always guaranteed.
Phase 2	Parallel Collection	Six intelligence streams dispatched simultaneously in one message. All subagent calls are schema-forced, eliminating the ~16% free-text JSON failure rate observed at scale. Results written to cache files in the session scratchpad.
Phase 3	Synthesis	One xhigh-effort agent reads merged.json by file path — no data tool access. Applies intelligence scoring (Critical → Background), meeting-first ordering, cross-source ♦ corroborated deduplication, and strict no-hallucination rules throughout.
Phase 4	QA Gate	HTML: qa-check.py validates anchors, attribution stamps, DOM containment, no light-mode flip, logo embed, junk patterns. Then one adversarial agent checks fabrication on dataFound=false accounts, struggle citations, and exec-claim traceability.
Phase 5	Deliver	Opens the output file, echoes the BLUF bullets and top 3 action items to chat, and reports any degraded or missing streams in Data Notes.

INTEGRATION REQUIREMENTS

Required — skill blocks in Phase -1 if any of these are missing or fail OAuth:

Salesforce (mcp__salesforce__)

PowerBI (mcp__powerBI__)

M365 connector (Enterprise built-in — not in claude mcp list)

Slack connector (Enterprise built-in — not in claude mcp list)

Optional — skill degrades gracefully:

LinkedIn (mcp__linkedin__)

→ WebSearch + public LinkedIn page fetch if absent/unauthed

SuccessFactors (mcp__successFactors__)

→ SF hierarchy + role/territory strategy only

Note: mcp__m365-copilot__* is *not* registered in claude-work and is not used by this skill. M365 email/calendar data comes from the claude.ai Enterprise connector above.

/exec-brief

Intelligence Streams · Output · Scoring · Guarantees

Version 1.0 · August 6, 2026

For: Dynatrace Executive Leaders

6 PARALLEL INTELLIGENCE STREAMS (PHASE 2)

SA-1 · M365 CALENDAR Calendar + Meeting Prioritization All customer-facing meetings in the window. Classifies by type (QBR, EBR, renewal, POC, escalation). Sets prep flags for exec-level attendees, escalation keywords, renewal proximity. Drives the entire brief's priority order.	SA-4 · SALESFORCE DEEP DIVE SF: Opps + Cases + Contracts + Activity 8 bulk SOQL queries in one agent over focus account IDs: open opps, closed won/lost, support cases (incl. escalated), active contracts, activity/tasks, account notes/feed, and POC/eval opps for prospects. Bulk pattern eliminates per-account agent failure class.
SA-2 · M365 OUTLOOK Email Intelligence Exec-keyed inbox scan, escalation sweep, and per-account passes (top 25 focus accounts). Calls read_resource on important items to get body text before summarizing. Surfaces action items, escalations, leadership signals, renewal negotiations.	SA-5 · POWERBI (HTTP, NOT AGENT) Consumption Data One HTTP call via pbi-extract.py returns the full ~500-account org grid in ~3.5s. Joined to book accounts by SFDC 18-char ID. Fields: health score, consumption rate (latest/6mo/Δ), ACV, risk level, forecasted % at contract end, open ATR, deployment type.
SA-3 · SLACK CONNECTOR Slack Intelligence Exec pass + per-account bare / risk / competitive query passes. Talking points only — max 3 bullets per account. Categories: struggles/escalations, competitive intel, AI-angle discussions, team/health mentions.	SA-6 · LINKEDIN + PUBLIC NEWS Public Intelligence Dynatrace news (earnings, launches, analyst) + per-company: WebSearch and LinkedIn (MCP or WebFetch fallback). Each company entry includes a mandatory DT alignment sentence — a cited public statement mapped to a specific DT capability play.

OUTPUT SECTIONS (9 ANCHORS IN HTML REPORT)

#exec-summary	BLUF bullets + 6-tile metric row (meetings · cases · renewals · pipeline · ACV · health)	#consumption	PowerBI: ACV · cons rate (latest/6mo/Δ) · forecasted % at end · end date · risk badge
#calendar	Meeting prep cards: attendees, type, 2–4 bullets (usage · struggle · consumption · the ask)	#opportunities	Pipeline by RD, top 10 deals, renewals calendar (90d/6mo/12mo), expansion signals
#key-accounts	Ranked by score — why flagged, owning AE + RD, headline numbers	#public-intelligence	DT news + per-company: dated headlines, LinkedIn, DT alignment sentence (cited → play)
#dt-usage	Deployment type, health, product footprint; POC status for prospects	#action-items	Ranked Critical→Low: account · action · urgency · horizon · owner (exec vs delegate)
#struggles	Source-cited only; ♦ cross-source corroborated outranks single-source		

FOCUS-ACCOUNT SCORING (PHASE 1E)

SIGNAL	POINTS
Upcoming customer meeting in lookahead	+100 (guaranteed)
Escalated / P1 open support case	+40
Open pipeline ≥ \$1M	+40
Open pipeline ≥ \$250K	+25
Matches FOCUS_CONTEXT keyword	+20
Prospect with open POC opportunity	+15
Any open case (cap +30 total)	+15
SF activity in lookback (cap +20)	+10

Cap: 50 focus accounts total. Meeting accounts are always guaranteed — the +100 score ensures they clear the cap unconditionally.

KEY DESIGN GUARANTEES

- ▶ **No hallucination — structural:** every finding requires a cited source. dataFound=false renders only the muted "no signals found" line. The adversarial QA agent hunts fabrication on empty-signal accounts.
- ▶ **Schema-forced outputs:** every subagent call uses { schema:... } — eliminates the ~16% free-text JSON failure rate (19/121 silent losses in a prior live run).
- ▶ **Meeting-first ordering:** the calendar drives the narrative. Accounts with upcoming meetings are guaranteed in the focus set regardless of the 50-account cap.
- ▶ **Bulk SOQL over per-account agents:** one WHERE AccountId IN (...) + code-side grouping covers all focus accounts. Removes the per-account agent parse-failure class entirely.
- ▶ **PowerBI = one HTTP call:** pbi-extract.py posts to the report's querydata endpoint — one ~3.5s call returns all ~500 org accounts. No browser scraping or agentic snapshot loop.
- ▶ **Cache-resume:** --resynthesize reruns synthesis from cached data without repulling. --refresh <stream> rebuilds one stream only. Good data is never clobbered.